Sibyl: Sales and Go-to-Market

Sibyl

14 October 2026

CONTENTS

Executive summary 5

Market problem: agents failing, and healthy agents doing forbidden things 7

- Agents now do work with consequences 7
- Problem one: the agent fails 7
- Problem two: the agent works and does something it should not 8
- Why existing controls do not close the gap 8

Ideal customer profile and buyer map 9

- Ideal customer profile 9
- Buyer map 11

Positioning and messaging 13

- One-liner 13
- Three pillars 13
- Messaging by persona 14
- Proof points tied to built capabilities 14
- Bring your own monitor agents 16
- Demo and real: links for prospects 18

Differentiation and competitive landscape 19

- Not just observability 19
- What the technical differentiation is 20
- The recovery boundary 20
- Competitive landscape 21

Where we do not compete 23

Objection handling 24

Use cases by function 28

Finance: payments and transfers 28

Customer refunds 28

Data exfiltration prevention 28

Insurance claims 29

Customer service and communications 29

Demo script 30

Design-partner program 34

Profile 34

What the partner commits to 34

What the partner gets 35

90-day pilot plan 35

Success criteria (agreed in writing at kickoff) 36

Finding the first partner on zero budget 36

Packaging and pricing hypotheses 38

Packages 38

Price shape 38

Sales process and qualification 40

Stages 40

MEDDICC-style qualification 41

Security review pack 41

Metrics 43
North star: governed autonomy 43
Supporting product metrics 43
Funnel metrics [estimates] 44
Pilot metrics 44
Launch readiness checklist 45
Product 45
Sales and marketing 45
Legal and company 46
Support 46
Appendix 47
What we say and what we never say 47
Claims to avoid 48
Glossary 49
Addendum: plug-and-play monitoring, drift detection and Sibyl Sentinel (built) 51
Plug-and-play monitoring 51
Native Anthropic and Gemini at the gateway 51
Model drift detection 51
Sibyl Sentinel 52
Action Gate status for sales 53

Executive summary

Sibyl is infrastructure that lets companies give AI agents more autonomy safely. It decides what an agent is allowed to do before the agent acts, keeps agents working when models and providers fail, and records every decision as evidence an auditor can verify offline.

Status as of 2026-10-14, release 1.0.0. This document separates three kinds of statement, and sales must keep them separate too:

- **Built.** Registry, a gateway that speaks the OpenAI, Anthropic and Gemini APIs, a Python SDK that instruments an agent in one line or none, OTLP ingest, explainable health, statistical leading indicators, tier-gated failover with approvals and undo, human handoff, kill switch, hash-chained audit, erasure, the console and the `sibyl-sim` provider doubles. Also built: the **Action Gate** (pre-execution allow, block or escalate on tool calls, from signed policy, at the gateway and in the SDK, with live console screens), **model drift detection** (watch-only statistics), **Sibyl Sentinel** (an advisory agent on the customer's own model key), **external monitor agents** (the customer's own Claude, OpenAI, Gemini, MCP or script-based monitors, advisory only), and two ways to run: a **demo** with dummy data and a **real deployment** for real agents, each with a Codespaces link.
- **Harness.** The Action Gate scenarios pass with no unauthorised action executed; 20 scenarios pass, including the Action Gate scenarios S-29 (approval and timeout re-run) and S-30, the silent model swap S-31, the outage reproduction S-32 and the new Sentinel scenario S-33; 12 do not yet pass, 11 of them waiting on a design ruling or an unbuilt feature and one a known coverage gap (B-10). Sentinel's judgement is not yet measured.
- **Research or hypothesis.** Learned prediction, cross-customer learning, pricing. None of these is claimed as fact.

The pitch in one paragraph. Monitoring catches an agent that breaks. It does not catch a healthy agent about to transfer $100,000 when it is authorised for $10,000, because every call in that sequence succeeds. That second failure class is what stops companies from giving agents real authority. Sibyl sits in the request path, evaluates each proposed

action against signed policy, holds risky actions for a human, and proves what happened. When the agent itself fails, Sibyl moves the work to an approved alternative or to a person, with undo and a kill switch.

Bring your own monitor. Prospects increasingly already have an agent they trust to review other agents. Sibyl lets them plug it in: their Claude, OpenAI or Gemini agent, an MCP assistant or a plain script reads the fleet and posts findings next to Sibyl's own Sentinel. It advises; Sibyl's signed, deterministic policy still decides. That turns "we already have an AI reviewer" from an objection into an integration.

Who we sell to first. Regulated or quasi-regulated companies of 500 to 5,000 people running 10 to 60 production agent workflows on two or more model providers, with a human exceptions team already in place. The Head of AI buys, the Platform Admin adopts, the CISO approves or vetoes, and risk and compliance co-sign.

What we are honest about. Sibyl is not a prediction engine today; its early warnings are statistics, measured in simulation only. Drift detection detects a silent model swap or a behaviour shift after it happens; it does not predict one. It has no data moat. It does not own the integrity of the customer's business transactions. Its advantage is integration in the request path, governance that survives an audit, and execution.

The commercial goal for the next two quarters is one design partner in a 90-day pilot with written success criteria, followed by a signed intent to convert. The budget for finding that partner is zero.

Market problem: agents failing, and healthy agents doing forbidden things

Agents now do work with consequences

Companies are moving agents from drafting text to taking actions: issuing refunds, moving money, approving claims, changing orders, emailing customers. Two different things go wrong, and in most organisations nobody owns either.

Problem one: the agent fails

Failure	What usually happens today
Provider outage or rate limit	Someone notices, edits configuration, redeploys
Silent quality loss (refusals, truncation, malformed tool calls)	Found days later through complaints
The provider serves a new model snapshot, or the agent’s behaviour shifts with no error	Nobody notices until outputs look wrong
Runaway cost or tool loops	Found on the invoice
Work abandoned mid-task	Lost, or done twice when someone retries

Observability tools help here, but they stop at the alert. Someone still has to decide what to do, at whatever hour it happens.

Problem two: the agent works and does something it should not

Example	Why monitoring misses it
Transfers $100,000 when authorised for $10,000	Every call succeeded; nothing looks unhealthy
Sends customer data to an external tool after a prompt injection	Normal latency, normal token counts
Issues 40 refunds in a minute, each individually valid	No single call breaks a rule

This is the larger risk. It is also the one that blocks adoption: a risk committee will not approve an autonomous refund agent because a dashboard shows it is fast. It approves when there is an independent limit the agent cannot exceed, a person on the high-value cases, and a record of every decision.

Why existing controls do not close the gap

- **Prompt instructions** (“never transfer more than $10,000”) are advice to the model, not enforcement. A prompt injection or a model change can override them.
- **Limits in the agent’s own code** are owned by the same team that ships the agent, change with every release, and are invisible to audit.
- **Limits in the downstream system** are the strongest control where they exist, but most internal APIs were built for trusted human operators, not for an automated caller that can make forty requests a minute.

Sibyl adds a control point that is independent of the model and of the agent’s code, sits before execution, and produces evidence.

Ideal customer profile and buyer map

Ideal customer profile

Attribute	Target	Why it matters
Industry	Financial services (lending, insurance, payments, broking), healthtech operations, BPO and customer-experience outsourcing, logistics with customer-facing agents	Real cost of a wrong or stuck action; existing exceptions teams; audits that ask about AI
Size	500 to 5,000 employees; AI platform team of 2 to 10	Large enough to have production agents and an audit function; small enough for the Head of AI to sign a pilot
Estate	10 to 60 production agent workflows; at least 3 tier-1 (money- or customer-facing); at least 2 model providers	Enough traffic for scores to mean something; tier 1 makes approvals relevant; two providers make failover real
Stack fit	OpenAI, Anthropic or Gemini clients that can point at a new base URL; Python for the tier-1 agents, ideally on LangChain/LangGraph, the OpenAI Agents SDK or CrewAI	Matches the built integration paths
Pain	A provider outage, rate-limit storm, model deprecation or near-miss action in the last 6 months that they can describe	The demo must land on a memory, not a hypothesis
Pressure	Operational-resilience or AI-governance expectations, or enterprise clients sending AI questionnaires	Gives the risk officer a reason to co-sponsor
Deployment	Willing to self-host in their own cloud account	Shortens their security review; keeps their data in their network

Disqualifiers. The estate is mainly closed SaaS agents with no call path to intercept. SOC 2 Type II is required before any pilot. 24/7 vendor support is required. They want Sibyl to replace their tracing or evaluation tool.

Buyer map

Role	Position in the deal	What they want	Objection they raise	Our answer
Head of AI	Economic buyer, usually the champion	Ship more autonomous workflows without more incidents; a defensible answer to the board	”Integration will cost my team a quarter.”	One base-URL change puts the gateway in front of an agent. In Python, one line (`sibyl.instrument()`) or none (`sibyl-run app.py`) instruments the clients and frameworks already in use. Pilot week one is observe-only.
Platform Admin	User; decides adoption	Zero-code onboarding, low overhead, no false failovers, undo	”Another proxy in my hot path.”	Measured overhead is about 0.2 ms at p99 in the CI benchmark. Policy is a local signed snapshot, so there is no control-plane call on the hot path. Each agent declares fail-open or fail-closed. The eject path is the same two-line change in reverse.
On-call Operator	User	One queue, a packet they can act on in minutes	”More alerts, more noise.”	Correlated failures collapse into one provider incident. Held actions arrive with the rule, the arguments and the reason. Handoff packets carry the original input, steps done and side effects.
Compliance Auditor	User, read-only	Evidence they can export without engineering help	”Logs I cannot read are not evidence.”	Each decision record shows the rule, the arithmetic and the approver. `sibyl audit verify` checks the exported chain offline.

Role	Position in the deal	What they want	Objection they raise	Our answer
CISO	Approver with veto	Independent enforcement, isolation, credential custody, kill switch	”You are a new vendor with no attestation.”	Self-hosted in their account. Row-level security tested on every table and role. Provider keys pass through and are not stored for the primary call. No SOC 2 yet; an external security assessment comes before the first paid tenant.
Risk or compliance officer	Approver or co-sponsor	Human oversight where it matters, retention, framework mapping	”Is this a compliance product?”	It supports human-oversight, inventory and record-keeping duties under NIST AI RMF, ISO/IEC 42001 and the EU AI Act. It does not certify compliance.
Procurement and legal	Process owner	Contracting entity, DPA, exit terms	”Who are we contracting with, and what happens when we leave?”	Pilot agreement with exit terms; full audit and payload export plus a crypto-shred erasure certificate on exit. (Entity and IP ownership must be settled before outreach; see the launch checklist.)

The central risk is the split: the Platform Admin decides whether Sibyl is adopted, the Head of AI whether it is bought, and the CISO whether it is allowed. Every deal needs all three.

Positioning and messaging

One-liner

Give your AI agents more autonomy, with guardrails, continuity and proof.

Category line: the control layer for the AI workforce.

Three pillars

Pillar	Claim	Status
Control	Sibyl decides what an agent may do before it acts: allow, block, or hold for a person. Policy is signed and enforced outside the agent.	Built: enforcement points, signed policy, the Action Gate, approvals, four-eyes and kill switch. Harness scenarios are merged; the full gate suite is not yet green.
Continuity	When a model or provider fails, work moves to an approved alternative or to a person, automatically where safe and with approval where not. Sibyl also notices when a provider silently changes the model.	Built. Drift detection is watch-only: it alerts a person and never moves traffic on its own.
Proof	Every decision, approval and override is in a hash-chained, signed log that verifies offline.	Built.

Messaging by persona

Persona	Message	Supporting line
Head of AI	Ship autonomous workflows without more incidents.	Raise an agent’s authority one policy at a time, with a person on the cases that matter.
CISO	An independent control layer the agent cannot talk its way past.	No model in the decision path; signed policy at two enforcement points; a kill switch that works when SSO is down.
Risk and compliance	Human oversight where it matters, evidence for every decision.	Approvals, holds, legal holds and erasure certificates, all exportable.
Platform Admin	Change one URL, or add one line, to start.	Works with the OpenAI, Anthropic and Gemini clients and with LangChain/LangGraph, the OpenAI Agents SDK and CrewAI. Undo every automated move; automation pauses for 30 minutes after an undo. Plug in your own monitor agent with one MCP entry or one key.
On-call Operator	One queue with the context to finish the job.	Packets with input, steps done, side effects and the next action.

Proof points tied to built capabilities

Use only these in external material. Each maps to something a prospect can see or test.

Proof point	Evidence
One-URL onboarding	Demo segment: an OpenAI client changes `base_url` and appears in the registry within seconds
One-line or zero-code monitoring	`sibyl.instrument()` or `sibyl-run app.py`; the console onboarding shows the snippet per client and framework
Native Anthropic and Gemini	Anthropic and Gemini SDKs change only base URL and key; the same kill switch, limits, routing and Action Gate apply. Fallback stays inside one protocol
Multi-agent visibility	A run tree shows each agent, step, model call, tool call and handoff; an unknown agent appears as discovered and is never routed until confirmed
Action control	Demo: an over-limit payout blocked before it runs; a mid-size payout held for two people; self-approval refused
Silent model swap detected	Demo: a new model snapshot raises `model_change`; a behaviour shift raises `silent_drift` after two drift verdicts; each report shows its per-feature numbers
Advisory explanations	Demo: Sentinel explains a failover and reviews a held payout, with model, tokens, cost and prompt hash on every finding
Bring your own monitor	Demo: the Monitors page lists a Claude reviewer over MCP, a webhook bridge to paging and a script; the held payout shows Sentinel's and the Claude monitor's reviews side by side, and the approver still decides. Live: `examples/monitors/` connects Claude, the OpenAI Agents SDK, Gemini, a webhook receiver and a rules script
Try it yourself	The public demo (https://sibyl-demo.onrender.com), the demo and real Codespaces links, or `make demo` and `make up` locally
Low overhead	Gateway CI benchmark, about 0.2 ms added at p99
Explainable decisions	Every health score shows `100 minus penalties` with each factor named
One incident, not seven	Correlated failures on a provider collapse into one provider incident

Proof point	Evidence
Governed failover	Automatic on detected failure; tier-1 backup-agent moves need approval; self-approval refused
Undo and hold	Undo restores the route and pauses automation for 30 minutes
Kill switch	Agent, provider or tenant scope; enforced at gateway and SDK; break-glass key works without SSO
Shared-dependency safety	Failover viability check refuses targets that share the failed provider
Human handoff	Structured packet with input marked untrusted, steps, side effects and cost
Tamper evidence	`sibyl audit verify` passes on the export and fails on a one-byte edit
Isolation	Second tenant stays silent through the outage; RLS tests per table and role
Early warning (simulation only)	In simulated ramps, warnings fired 5 to 39 minutes before failure
Nightly gate report	Scenario harness across seeds, with CRITICAL gates at 100 percent (scenario runner still in the harness lane)

The harness proof point for the Action Gate (“zero unauthorised actions executed across the harness scenarios”) is added only when the A-5 gate is green. Drift detection and Sentinel have no harness numbers yet either; do not quote any. External monitors are never measured by Sibyl; there is no quality figure to quote for them.

Bring your own monitor agents

The positioning. Sibyl is an open monitoring ecosystem with a closed enforcement core. Any agent the customer trusts can watch the fleet; only Sibyl’s signed, deterministic policy can act.

What the prospect has	How it plugs in	What it can do	What it cannot do
A Claude assistant (Desktop, Claude Code, Agent SDK)	`sibyl-mcp`, one MCP entry	Read agents, health, runs, drift, alerts, incidents and held actions; post findings and held-action advice	Allow, block, approve, fail over, change policy
An OpenAI Agents SDK or Gemini app	`sibyl-mcp`, or the SDK's `sibyl.monitor` client	The same	The same
Paging, chat or ticketing tools	Signed webhooks for alert, incident, held-action, drift and failed-run events	Receive events and post findings back	The same
A rules script or cron job	REST API with a monitor key	The same	The same

What to say. "Bring the reviewer you already trust. Register it in the console, give it a key that can only read metadata and post advice, and its findings appear next to Sentinel's, on the very held action or incident they concern. It never decides: the Action Gate, the approver and the failover rules do."

Why it matters to each buyer.

- **Head of AI:** keeps their own model and prompt choices; Sibyl does not force a reviewer on them.
- **CISO:** the monitor's key holds two scopes, sees only its watched agent groups, reads metadata only, and cannot reach any route that acts. A compromised or prompt-injected monitor costs a bad recommendation, not a transfer.
- **Risk and compliance:** every registration, finding, rating and failed delivery is in the audit chain.
- **On-call Operator:** events arrive in the tools they already watch, signed and retried.

Limits to state. Monitors read metadata only (no prompts, outputs or held-action arguments). Monitor keys get webhooks or polling, not a live stream. The only run event is `run.failed`. A monitor can see other monitors' findings within its groups. A paused

monitor's webhooks wait until it is resumed. Sibyl does not measure a monitor's quality. `sibyl-mcp` and `sibyl-sdk` install from the repository, not PyPI, for now.

Demo and real: links for prospects

	Demo (dummy data)	Real deployment (their agents)
Link	Open in Codespaces (demo)	Open in Codespaces (real)
Locally	`make demo` (Node only)	`make up` (Docker with Compose)
Time to first screen	A couple of minutes	About 10 minutes
Use it for	Discovery calls and first demos: every screen, a scripted outage, held actions, drift, Sentinel and monitors	Technical validation (stage 3): the Platform Admin connects one real agent, approves a first policy and registers a monitor

Access. The repository is public. Send the public demo (https://sibyl-demo.onrender.com: dummy data, shared by every visitor, no sign-up) for a first look, and the Codespaces links for a private copy of the demo or a real deployment. The real deployment in a codespace is for evaluation, not production; production is self-hosted in the prospect's own environment.

Differentiation and competitive landscape

Not just observability

Observability tools record what happened. Sibyl acts, under governance, and proves it.

Capability	Observability tools	Sibyl
See calls, latency, errors, cost	Yes	Yes (not our point of difference)
Decide what an agent may do before it acts	No	Yes: signed policy at gateway and SDK, argument-level rules and cumulative caps
Notice a silent model swap or behaviour shift	Sometimes, as a chart	A `model_change` or `silent_drift` alert with per-feature statistics; watch-only by design
Let your own AI reviewer watch the fleet	Usually their own evaluator only	Any Claude, OpenAI, Gemini, MCP or script monitor, with a scoped key, signed webhooks and findings beside the decision they concern; advisory only
Act on failure	Alert only	Failover by risk tier, approvals, undo, hold, kill switch
Continue unfinished work	No	Handoff packet to a human queue
Tamper-evident evidence	Not a design goal	Hash chain, signed anchors, offline verification

What the technical differentiation is

None of these is novel alone. The combination, enforced in the path and evidenced end to end, is the product.

1. **In-path enforcement.** Two enforcement points: a gateway (OpenAI, Anthropic and Gemini protocols) that sees each model response, including proposed tool calls, before the agent acts on it; and an SDK check that runs immediately before each tool executes, including tools run by LangChain, the OpenAI Agents SDK and CrewAI.
2. **Signed policy.** Enforcement points act only on Ed25519-signed policy snapshots. A forged or stale snapshot is refused; the kill state is sticky.
3. **Approvals.** Risky actions are held for a person. Four-eyes for tier 1. The requester cannot approve their own request. Timeouts default to deny for held actions.
4. **Undo.** Every automated move can be reversed, and reversal pauses automation.
5. **Evidence.** Every decision carries its rule, its arithmetic and its approver in a hash-chained log.
6. **No LLM in the decision path.** Detection is deterministic rules and change-point statistics; drift detection is PSI and Jensen–Shannon against a frozen baseline. The Action Gate evaluator is a pure function over a JSON expression tree, with no `eval`. The agent cannot persuade it, because nothing in the decision reads the agent's prose. Sentinel and the customer's own monitor agents use models, but only to advise a person.

The recovery boundary

Sales must state this boundary unprompted, because buyers who discover it later lose trust.

Sibyl owns	The customer’s infrastructure owns
Which model, provider or agent a call goes to	Business transaction integrity
Stopping traffic (kill switch, quarantine)	Idempotency of their tools and APIs
Handing unfinished work to a person, with context	Re-running or compensating side effects
Evidence of every decision	Their own data stores and backups

- **Shared upstream dependencies.** If every alternative depends on the same failed API, switching achieves nothing. Sibyl records dependencies, refuses failover targets that share the failure, and sends the work to a person or a degraded mode instead.
- **Duplicate side effects.** A task whose recorded steps include an external side effect is never resumed automatically. A person decides, with the side-effect ledger in front of them. Automatic resume applies only to steps the agent declared idempotent. (Built: backlog item A-4.)
- **Defence in depth.** For the highest-risk tools, we recommend the downstream system keeps its own limits as well. Sibyl’s gate is independent of the agent, not a replacement for the bank’s own controls.

Competitive landscape

Market facts below come from public sources summarised in October 2026 and have not been tested by us. Items marked **[unverified]** must not appear in external material without fresh confirmation.

Category	Examples	Where they stop	Where Sibyl starts
LLM observability and evals	Langfuse (acquired by ClickHouse, Jan 2026), Arize Phoenix (ELv2), LangSmith, Datadog LLM Observability	Not in the request path; no routing, approvals, handoff or decision evidence	Control and continuity in the path; we accept their OTLP and sit next to them
Instrumentation	OpenLLMetry, OpenInference	Emit spans; make no decisions	An input to Sibyl, not a competitor
Gateways and routers	LiteLLM, Portkey, Kong AI Gateway, Helicone (reported in maintenance mode after acquisition **[unverified]**)	Request-level, config-driven fallbacks; no tier gating, human queue or tamper-evident decision record **[unverified: current approval and guardrail features in commercial tiers]**	Governed, explainable, undoable decisions; handoff; evidence. LiteLLM is the closest overlap
Content guardrails	Prompt and output filtering products **[unverified: feature scope]**	Typically judge text content, often with a model	Argument-level policy on actions (amounts, destinations, cumulative caps), deterministic and signed
Enterprise AI governance suites	ServiceNow AI Control Tower, Credo AI, Holistic AI, IBM watsonx.governance; Microsoft Agent 365 **[unverified: GA status]**	Strong on registers, assessments and framework content; mostly not in the request path **[unverified: ServiceNow states it can terminate**	Evidence generated by operations, not assessments; vendor-neutral; self-hostable

Category	Examples	Where they stop	Where Sibyl starts
		agents beyond permissions]	

Market read: in 2026 several observability players were absorbed into larger platforms, and governance platforms are buying telemetry. A standalone observability pitch would be a race to the bottom.

Where we do not compete

We do not build	Instead
Trace waterfalls, span search, session replay	Show the slice that explains a decision; link out by trace id
Evals, datasets, LLM-as-judge, prompt playgrounds	Consume eval results as a signal
Prompt management	Record `prompt_version` as an attribute
Breadth across 100+ providers, semantic caching	Support the dialects that matter; chain behind LiteLLM where breadth is needed
Line-item cost dashboards	Budget as a health factor and limit
Closed SaaS agents with no interceptable call path	Disqualify; revisit with platform integrations later

Tell prospects to keep their observability tool. A pilot that requires ripping one out will stall in week one.

Objection handling

#	Objection	Answer
1	”Isn’t this just monitoring?”	Monitoring tells you what happened. Sibyl decides what is allowed to happen: it sits in the request path, holds or blocks actions against signed policy, moves work when a provider fails, and lets a person undo it. If we only showed dashboards, you should keep your current tool.
2	”An agent watching agents isn’t a breakthrough.”	Agreed, and that is not what this is. No LLM sits in the detection or decision path. Rules and the action evaluator are deterministic and explainable, and policy is signed. The agent cannot argue with a predicate like `amount <= 10000`. Sentinel, the optional model-based reviewer, only advises a person.
3	”You can’t recover my business transaction.”	Correct, and we say so. Your systems own transaction integrity, idempotency and compensation. Sibyl owns routing, stopping traffic, handing work to a person, and the evidence. A task with recorded external side effects is never resumed automatically; a person decides with the side-effect ledger.
4	”Where’s your moat?”	We do not claim a data moat. Our advantage is integration in the path (one URL change), governance that holds up in an audit, and execution: self-hosted, one-command deploy, safe defaults. Candidates for a future moat are per-tenant verdict-labelled incident data, the scenario library and policy templates, each to be proven by measured lift.
5	”Can you predict failures?”	Today we raise statistical early warnings (trend and change-point detection, rate-limit and budget proximity). In simulated ramps they fired 5 to 39 minutes ahead. That is simulation, not your traffic. A learned model ships only if it beats the rules on your own labelled alerts, and you will see the numbers either way.

#	Objection	Answer
6	”How do you isolate our data?”	Partner one self-hosts in their own cloud. Row-level security on every tenant table, tested per table and role. Per-tenant envelope encryption. No pooling of signals or models across tenants (a standing decision). Capture can be full, redacted or metadata only, and secrets are scrubbed at three points.
7	”Another proxy in my hot path is a single point of failure.”	Measured overhead is about 0.2 ms at p99. The gateway runs on a signed local snapshot with no control-plane call per request. Each agent is fail-open or fail-closed by declaration. If you want out, restore your original base URL.
8	”Our LLM gateway already does fallbacks.”	Keep it if you like; Sibyl can sit in front of it. Static fallback chains do not know risk tier, do not wait for approval on tier 1, cannot be undone with a hold, do not refuse a target that shares the failed dependency, and do not hand unfinished work to a person.
9	”We already have guardrails.”	Content guardrails judge text. The Action Gate judges actions: tool, arguments, destination, cumulative totals, data class. It is deterministic and signed, and it can hold an action for a human instead of only refusing it. They are complementary.
10	”What stops the agent bypassing Sibyl?”	Sibyl enforces where traffic flows through it. Tier-1 agents must use agent-bound keys, and auto-discovery registers unknown agents seen in traffic or as handoff targets; they are never routed until an admin confirms them. Code that calls tools without the SDK or gateway is outside coverage; that is why our north-star metric measures the share of actions under enforcement. For critical tools, keep downstream limits too.
11	”You have no SOC 2 and one operator.”	True. That is why partner one self-hosts, support is business hours stated in writing, fail-open and the eject runbook cover out-of-hours, and an external security assessment comes before any paid contract. SOC 2 readiness is a Phase 3 item.
12	”Approvals will slow everything down.”	Approvals apply only where policy says so: above a threshold, on sensitive tools, or on tier 1. Everything else is automatic. Timeouts are explicit, and held actions default to deny so nothing executes unseen. We measure approval latency in the pilot.

#	Objection	Answer
13	”Is this an EU AI Act compliance product?”	It supports human-oversight, inventory and record-keeping duties and produces evidence for them. It does not certify compliance. We lead with operational resilience; high-risk deadlines have reportedly moved to late 2027 **[unverified: Official Journal publication]**.
14	”What if Sibyl makes a wrong decision?”	Every automated move can be undone in one click, which pauses automation for 30 minutes. Two undone moves in seven days switch that agent to approval-only. A single weak signal never triggers automation. Every decision shows its arithmetic.
15	”Will this lock us in?”	The SDK is Apache-2.0. OTLP ingest means existing instrumentation keeps working. On exit you get full audit and payload export and a crypto-shred erasure certificate.
16	”Is the Action Gate real?”	Yes. The evaluator, enforcement at the gateway and SDK, held actions with four-eyes, self-approval refusal, cumulative caps and the console screens are built and tested. What is still in progress is the harness result across over-limit, cumulative-cap, exfiltration and injection scenarios; we will show you that number when it is green, not before.
17	”How much work is it to instrument our agents?”	For Python, one line (`sibyl.instrument()`) or none (`sibyl-run app.py`). It picks up the OpenAI, Anthropic and Google Gen AI clients and LangChain/LangGraph, the OpenAI Agents SDK and CrewAI. Libraries you do not use are skipped. If configuration is wrong, your app still starts and tells you why. Google ADK and AutoGen connectors are planned; OTLP works today.
18	”Can you tell us when our provider quietly changes the model?”	Yes, after it happens. Sibyl compares the model the provider says it served, and its fingerprint where the provider sends one, with what each agent was served before. It also compares output length, latency, refusals, tool use and errors against a frozen baseline, with a noise floor so small samples do not raise alarms. You get an alert with the numbers. We do not predict drift, and drift alone never moves traffic.

#	Objection	Answer
19	”Your Sentinel is an LLM. Why trust it?”	Do not trust it to decide, and it cannot. It explains incidents, gives the approver a second opinion on a held action, flags purpose drift and summarises the fleet. It cannot allow, block, approve, fail over or change policy. It runs on your own key, within a daily budget you set, sends no prompts or outputs for metadata-only agents, and records model, tokens, cost and a prompt hash on every finding. We have not yet measured its quality on our harness.
20	”We use Claude and Gemini, not OpenAI.”	The gateway speaks the native Anthropic Messages and Gemini APIs, streaming included, with the same controls. Your SDK changes only its base URL and key. Fallback stays within one protocol: an Anthropic-protocol agent fails over to Anthropic-protocol targets. Sibyl does not translate between protocols.
21	”We already have our own agent reviewing our agents.”	Good: plug it in. Register it as a monitor and connect it through the Sibyl MCP server, signed webhooks or the API. Its findings appear next to Sentinel’s, on the held action or incident they concern. It stays advisory: its key reads metadata in the groups you choose and posts findings, and cannot reach any route that allows, blocks, approves or fails over.
22	”If our monitor agent gets prompt-injected, what can it do?”	Write a wrong finding, which a person reads and can rate as wrong. It cannot act: two scopes only, watched groups only, metadata only, no enforcement route, rate-limited, and every finding audited. Revoking the monitor revokes its keys and disables its webhooks at once.
23	”Can our monitor read the prompts and outputs?”	Not today. Monitor reads are metadata only, whatever your capture mode. For semantic review of content, Sentinel reads content within the capture mode; a content path for external monitors is on the roadmap, not built.
24	”Can we try it without a call?”	Yes: the public demo at https://sibyl-demo.onrender.com (dummy data, every screen), the demo Codespaces link or `make demo`; the real-deployment link or `make up` to connect your own agents.

Use cases by function

Each use case states the built continuity value and the Action Gate angle. Policy examples use the Action Gate rule model: tool, argument predicates, cumulative caps, data class and effect.

Finance: payments and transfers

- **Workflow.** An agent prepares and executes vendor payments or treasury transfers.
- **Action Gate.** `payments.transfer`: allow when `amount <= 10000` and `destination in approved_beneficiaries`; escalate with dual control between 10,000 and 25,000; block above 25,000 or to any unlisted destination. Daily cumulative cap per agent with escalate on breach.
- **Continuity.** Tier 1. A backup-agent move needs approval. A transfer step with a recorded side effect is never auto-resumed.
- **Evidence.** Each held transfer shows the rule, the amount, both approvers and timestamps.

Customer refunds

- **Workflow.** A support agent issues refunds and goodwill credits.
- **Action Gate.** Per-refund cap; cumulative caps per customer per day (sum) and per agent per minute (count, e.g. 20 per 60 seconds). The 40-refunds-in-a-minute pattern escalates at the 21st.
- **Continuity.** Tier 2. Provider failures move traffic automatically; unfinishable cases go to the exceptions team as packets.

Data exfiltration prevention

- **Workflow.** Any agent with access to customer records and external tools (email, webhooks, file upload).

- **Action Gate.** Block when an argument scrubs as a secret or matches PII detectors and the tool is marked external. Tool allowlists block tools the agent never declared. An injection-driven call to a new destination is blocked by the destination predicate regardless of how the model was persuaded.
- **Continuity.** Kill switch for a compromised agent in seconds.

Insurance claims

- **Workflow.** An agent triages claims and proposes settlements.
- **Action Gate.** Allow settlements below a threshold for low-risk claim classes; escalate above it or when the fraud flag is set; block payment to a bank account changed in the last N days.
- **Continuity.** Claims in flight during an outage go to adjusters with the full packet, so nothing is lost or paid twice.
- **Evidence.** Human-oversight records for each escalated settlement, exportable for internal audit.

Customer service and communications

- **Workflow.** Agents answer, update orders, and send outbound messages.
- **Action Gate.** Escalate outbound messages to a segment above a size; block order cancellations above a value; allow read-only tools freely.
- **Continuity.** Rate-limit storms move traffic to the approved alternate model; quality loss (refusals, truncation, malformed tool calls) lowers the health score with the cause named.

Demo script

Fifteen minutes, three role-based personas: Platform Admin, On-call Operator, Compliance Auditor. The demo runs on the console's demo-data mode (a purple **Demo data** chip is on every screen): the synthetic tenant `meridian` (Meridian Logistics) plays a scripted story in which Anthropic degrades and then goes down, while the Action Gate, run tree, drift and Sentinel beats play alongside. `northstar` is the isolation proof and `sandbox` is empty. Real providers are off. Three browser profiles are signed in. Say once, at the start, that this is demo data shaped exactly like the live API.

What to say about the harness. The Action Gate is built and works as shown. Do not claim a harness result for it until the A-5 gate is green. Do not quote any accuracy number for drift detection or Sentinel.

Time	Segment	Persona	What the audience sees (console demo beat)	Proof point
0:00 to 1:00	The hook	Platform Admin	Fleet of agents, each with owner, tier and continuity plan. One tier-1 PII agent is parked on hold because its approval timed out and it was not moved without a human.	Governance by default
1:00 to 2:30	Onboard from nothing	Platform Admin	Onboarding: capture policy first, then the first key. The snippet tabs show one line or zero code for OpenAI, Anthropic, Gemini, LangChain/LangGraph, the OpenAI Agents SDK and CrewAI. A gateway call registers a new agent within seconds.	One line, or one URL
2:30 to 4:00	Blocked payout in a claims run	Platform Admin	A ₹2,50,000 payout (limit ₹1,00,000) is blocked by the Action Gate. Open the claims run tree: each step, the model calls with the served model, and the payout tool call blocked in the run. A second beat shows personal data to an outside email blocked.	Control before execution; multi-agent visibility
4:00 to 5:30	Outage and failover	Platform Admin	Anthropic latency climbs, then every call returns 503. Error-rate alerts on 7 agents collapse into one provider incident. Tiers 1, 2 and 3 fail over to OpenAI within their budgets. An agent’s score shows its arithmetic.	Explainable, automatic, one incident
5:30 to 7:30	Handoff to an unknown agent	Platform Admin, On-call Operator	The claims run hands off to a payout agent Sibyl has never seen. It appears under Discovered agents as unconfirmed and not routed. The run’s ₹75,000 transfer is held for two people. It was proposed acting for	Discovery; four-eyes; no self-approval

Time	Segment	Persona	What the audience sees (console demo beat)	Proof point
			the Platform Admin, so the Platform Admin's own approval is refused (`self_approval`); the On-call Operator approves with step-up, and four-eyes needs a second approver.	
7:30 to 9:00	Sentinel advises	On-call Operator	Sentinel explains the claims-triage failover in plain language, and reviews the ₹75,000 hold: advice `needs_more_info` , reasons, risks, and provenance (model, tokens, cost, prompt hash). The banner says Sentinel never decides.	Advisory explanation, on the customer's key
9:00 to 10:30	Silent model change and drift	Platform Admin	The sales copilot is served a new gpt-5 snapshot: a `model_change` alert. The route optimizer's tool sequences and latency shift: after two drift verdicts in a row, a `silent_drift` alert. Open the agent's drift panel: each feature's value, noise floor and threshold. Sentinel then reads sampled runs and finds no sign of injection. Both alerts are watch-only.	Detects swaps and shifts; explains, never acts alone
10:30 to 11:30	Operator queue and refunds	On-call Operator	A refund burst hits the 5-a-minute cap and the sixth is held. A handoff packet arrives with input marked untrusted, steps done and side effects; it was not auto-resumed. Closing an alert needs a verdict.	Cumulative caps; human handoff
11:30 to 12:30	Undo and kill	On-call Operator	Undo one failover with a reason; the banner shows automation paused. Pull the kill switch on another agent with step-up.	The human always wins

Time	Segment	Persona	What the audience sees (console demo beat)	Proof point
12:30 to 14:00	The auditor	Compliance Auditor	Audit view filtered to Action Gate decisions: the blocked payout, the held transfer, the refused self-approval, the kill, each with its rule. Export; `sibyl audit verify` reports the chain intact; a one-byte edit fails.	Tamper-evident, verifiable offline
14:00 to 15:00	Isolation and close	Compliance Auditor	Switch to `northstar`: zero alerts and zero transitions throughout. Close on the recovery boundary and what is still in the harness lane.	Isolation; honesty

Optional Q&A beats. Rate-limit early warning on a tier-1 agent shown as an approval card (state clearly: statistical, simulation-measured). Sentinel setup: provider, write-only key, daily budget, sampling rate. **Bring your own monitor** (2 minutes, for prospects who already run a reviewer agent): open **Monitors**, show the Claude reviewer over MCP, the webhook bridge with its delivery status and the paused script; register a new monitor and point at the key shown once with its two scopes; then open the ₹75,000 hold and show Sentinel's and the Claude monitor's reviews side by side, with the Decide card untouched. Tenant erasure in `sandbox` producing a certificate. A live terminal run against the compose stack for technical audiences.

Rehearsal rule. Three dress rehearsals from a cold start; two consecutive runs with no fallback before any external demo; a recorded, chaptered fallback video; never demo a build that failed that night's harness.

Design-partner program

Profile

The ideal customer profile in chapter 3, scored as a checklist. A prospect missing “pain” or “stack fit” is not a design partner, however enthusiastic. Add one Action Gate criterion: at least one workflow where the partner would grant more authority if an independent limit and an approval existed.

What the partner commits to

Commitment	Ask
Sponsor	Named Head of AI at kickoff, day 45 and day 90 reviews
Engineering	One Platform Admin equivalent, about 4 hours a week
Operators	Two people who already handle exceptions work the Sibyl queue and held actions
Risk and audit	One reviewer who reads an audit export at day 45 and day 90 and states in writing whether it is usable as evidence
Traffic	At least 3 agents (1 tier 1) via the gateway; target 60 percent of LLM calls by day 60
Labels	A verdict on every alert and every blocked action
Drills	One game-day a month with operators present
Feedback	30 minutes every two weeks; access to anonymised incident notes
Reference	Named or anonymised case study
Paperwork	Mutual NDA, DPA covering capture mode and retention, pilot agreement with exit terms

What the partner gets

Free use for 90 days; founding-partner pricing **[hypothesis]**; direct founder access with business-hours support; a say in sequencing two roadmap items; self-hosting in their account with a full exit path; their own per-tenant model in shadow once verdict volume allows.

90-day pilot plan

Weeks	Phase	Activities	Exit check
0	Kickoff	Success criteria signed; onboarding workshop: capture mode, tiers, continuity plans, fail modes, agent-bound keys for tier 1; eject runbook rehearsed	Criteria and eject tested
1 to 2	Observe	3+ agents through the gateway; no routing changes; action policies in simulate mode only	Zero Sibyl-caused incidents; overhead within budget
3 to 6	Govern continuity	Failover on for tiers 2 and 3; tier 1 under approvals; first drill	First real or drilled failover per tier-1 agent
5 to 8	Govern actions	Action policies enforced on one or two sensitive tools; escalations to operators; weekly review of blocks and holds	Day 45 evidence review; zero false blocks unexplained
9 to 12	Expand and prove	Coverage toward 80 percent; second drill; willingness-to-pay conversation at day 60	Day 90 review; conversion decision

Success criteria (agreed in writing at kickoff)

#	Criterion	Threshold
1	Governed autonomy (share of actions under approved policy)	At least 60 percent at day 60, 80 percent at day 90
2	Sibyl-caused production incidents	0
3	Gateway overhead p95 on their traffic	At most 15 ms
4	Unauthorised actions executed on governed tools	0
5	Continuity proven	At least 1 failover per tier-1 agent; at least 90 percent of tier-1 tasks completed despite failure
6	False automated failovers on tier 1	0
7	Alert precision from their verdicts	At least 70 percent; at least 90 percent verdicted
8	Operator usefulness	Median packet-to-complete at most 10 minutes
9	Evidence accepted	Written statement from their risk or audit owner
10	Retention and commercial signal	Sponsor “very disappointed” without Sibyl; signed intent at a price band, or a named reason why not

A pilot that meets 1 to 9 but fails 10 is a product success and a positioning failure, and we record it as such.

Finding the first partner on zero budget

Precondition: entity and IP ownership confirmed before any external conversation (an open founder item).

Channel	Action	Expected yield [estimate]
Founder network	30 direct messages to Heads of AI in regulated mid-market firms, asking for 20 minutes on “how do you limit what your agents can do, and what happens when a provider fails?”	10 conversations
Proof content	A 3-minute demo video and a write-up using harness numbers; LinkedIn, Show HN, MLOps Community	3 to 5 inbound
Incident timing	Same-week note with a reproducible scenario after a public provider outage	Interest spikes
Open-source SDK	Publish `sibyl-sdk` (Apache-2.0); examples for LangGraph and OpenAI Agents SDK	Engineer-led leads
Risk circles	A 30-minute talk on “evidence for AI human oversight” to ISACA or risk-practitioner chapters	1 to 2 risk-side introductions

Funnel target **[estimate]**: 40 touches, 15 discovery calls, 6 demos, 3 qualified, 1 signed pilot within 8 weeks of the Phase 1.5 exit.

Packaging and pricing hypotheses

Everything in this chapter is unvalidated. No price is quoted to a prospect before at least three pricing conversations with qualified prospects.

Packages

Package	For	Includes
Developer (free)	Evaluation	SDK, gateway, registry, health, single tenant
Team	First production agents	Action Gate, failover, handoff queue, audit export
Business	Several teams and providers	Approvals and four-eyes, leading indicators, legal holds, erasure certificates, SSO
Enterprise	Regulated firms	Compliance packs, long audit retention, residency, support SLA

Price shape

- **Platform fee plus a per-governed-agent charge, weighted by tier** (hypothesis: tier 1 = 3 units, tier 2 = 2, tier 3 = 1), with a fair-use event band. Weighting must never reward under-declaring tier.
- **Indicative ranges [unvalidated]:** Team about USD 1,000 to 2,000 a month platform fee plus USD 40 to 80 per agent-unit a month; Business about USD 4,000 to 6,000 a month platform fee with a lower unit rate; Enterprise custom, floor about USD 75,000 a year.
- **Add-ons:** compliance packs per framework per year; audit retention beyond 400 days.
- **Founding partner [hypothesis]:** 50 percent off Team or Business list for 24 months if converted within 60 days of pilot exit.
- **Never paywalled:** the SDK, OTLP ingest and the kill switch.

Rationale: priced against the cost of one serious incident or one unauthorised payment, and a fraction of a compliance analyst's time, not against per-trace observability pricing. Validate with Van Westendorp questions at day 60 of the pilot.

Sales process and qualification

Stages

Stage	Exit criteria	Owner artefact
0. Target	Matches the ICP on paper	Account note
1. Discovery	Pain described with a specific incident or blocked workflow; stack fit confirmed	Discovery notes
2. Demo	12-minute demo to Head of AI and Platform Admin; next step booked	Demo recording
3. Technical validation	Platform Admin runs the real deployment (`make up`, or the real Codespaces link) in their sandbox; one agent through the gateway or the SDK; optionally their own monitor agent registered	Validation checklist
4. Security review	CISO receives the security pack; no blocking finding	Review log
5. Pilot agreement	Signed success criteria, DPA, exit terms	Pilot agreement
6. Pilot	90-day plan executed	Weekly status
7. Conversion	Intent or contract at a stated band, or a documented no	Exit review

MEDDICC-style qualification

Element	What we need to know	Disqualify if
Metrics	Which number improves: incidents, tasks completed, workflows granted more authority, audit effort	No measurable outcome
Economic buyer	Head of AI (or equivalent) with authority to sign an unpaid pilot	Buyer unknown after stage 2
Decision criteria	Overhead, isolation, self-hosting, evidence quality, action control	Criteria require SOC 2 Type II now
Decision process	Who approves a pilot: security, legal, procurement steps and timing	Process exceeds 90 days for an unpaid pilot
Identify pain	A described outage, near-miss action, or workflow blocked by risk	Pain is hypothetical
Champion	Someone who will run the internal review and defend the pilot	No champion by stage 3
Competition	What they use now (observability, gateway, governance suite) and what they would keep	They want a tracing replacement

Security review pack

Contents, with honest status:

Item	Status
Security overview (architecture, data flow, encryption, isolation, audit chain)	To write before stage 4
Data handling: capture modes, three-point secrets scrubbing with 23 detectors, per-tenant AES-256-GCM envelope encryption, retention	Built
Credential custody: passthrough for primary calls, vaulted fallback keys, hashed API keys	Built
Isolation evidence: RLS tests per table and role; second-tenant demo proof	Built
Signed policy and kill switch design, including break-glass key	Built
Harness gate report (latest nightly)	Built
SBOM and licence inventory (permissive dependencies only)	Built in CI
Known limits page (accepted residuals, business-hours support, declared residency)	To publish
Incident response and breach notification procedure	To write with counsel
DPA template and sub-processor list (none for self-hosted)	To write with counsel
External security assessment report	**Not done**; required before the first paid tenant
SOC 2	**Not started**; Phase 3

Metrics

North star: governed autonomy

Governed autonomy is the share of a tenant's production agent actions that pass through Sibyl's enforcement points under a declared, approved policy.

It moves only when Sibyl is adopted (actions flow through it), governed (policies exist and are approved) and trusted (teams route more authority through it). It refines the earlier "verified governed coverage" measure, which counted calls from agents with an owner, a tier, a risk class and a recently proven fallback; those conditions remain as supporting checks. Sibyl cannot see actions that bypass it, so discovering unknown agents can lower the number. We report that drop as progress, not failure.

Counter-metrics (any one red blocks a claim of progress): Sibyl-caused incidents, false automated failovers, false blocks, gateway overhead p95, unverdicted backlog.

Supporting product metrics

Metric	Why
Unauthorised actions blocked or escalated	Direct value of the Action Gate
Unauthorised actions executed (harness and pilot)	Must be zero
Held-action decision time, median	Approvals must not stall work
Time to detect	Continuity
Failovers within recovery target	Continuity
False-failover rate by tier	Trust
Tasks completed despite agent failure	Business continuity
Verdict coverage	The only source of per-tenant learning

Funnel metrics [estimates]

Stage	Target per cycle
Outreach touches	40
Discovery conversations	15
Demos	6
Qualified (ICP scorecard)	3
Signed pilots	1 within 8 weeks of Phase 1.5 exit
Pilot to intent to convert	1 of 1 for partner one; measured thereafter

Pilot metrics

The ten success criteria in chapter 9, reviewed every two weeks with the partner, plus founder support hours (target at most 6 a week by day 60).

Launch readiness checklist

Product

- [x] Action Gate A-1 to A-4 and A-6 merged, with the console on the live contract
- [ ] A-5 harness gate green (zero unauthorised actions executed)
- [x] A-4 enforced: side-effecting tasks never auto-resumed, with tests
- [ ] Silent model swap scenario in the harness; outage detection gap closed
- [ ] Sentinel quality measured on the harness before it is offered as on by default
- [ ] Full nightly harness green on five consecutive nights
- [ ] Backup restore drill passed, including re-application of erasures
- [ ] Gateway overhead benchmark within budget with scrubbing and encryption on
- [x] Self-host guide (`docs/setup/SETUP.md`) and the real deployment (`make up`) with OIDC and TLS options
- [x] External monitor agents: registration, monitor keys, `sibyl-mcp`, signed webhooks, examples
- [x] Release 1.0.0: GitHub Actions CI, release images on GHCR, demo and real Codespaces links
- [ ] Helm chart for partner one; eject runbook written
- [x] Public demo (https://sibyl-demo.onrender.com) and public repository with Codespaces links

Sales and marketing

- [ ] Demo rehearsed three times; two consecutive clean runs; chaptered fallback video
- [ ] Security review pack complete (chapter 11), including Known limits
- [ ] Claims review: every external sentence checked against chapter 14 and the “What we say” table
- [ ] 3-minute demo video and technical write-up
- [ ] ICP scorecard and discovery script
- [ ] `sibyl-sdk` and `sibyl-mcp` published under Apache-2.0 (package name reservation is an open founder item)

Legal and company

- ☐ Contracting entity and IP ownership confirmed (open founder item) before outreach
- ☐ Pilot agreement, DPA and privacy policy reviewed by counsel
- ☐ External security assessment scheduled; closed with no open Critical or High before the first invoice

Support

- ☐ Business hours and timezone published; severity table with first-response targets
- ☐ Shared support alias; status page on separate infrastructure
- ☐ Breach notification procedure agreed with counsel

Appendix

What we say and what we never say

What we say	What we never say
Sibyl controls what agents may do before they act.	”Sibyl makes your agents safe.”
The Action Gate is built and blocks or holds actions before they run. In our harness runs, no unauthorised action executed in the Action Gate scenarios.	”Zero unauthorised actions” or “prevents all unauthorised actions” (before A-5 passes, and never as an absolute).
Early warnings are statistics; in simulation they fired 5 to 39 minutes ahead.	”Predictive intelligence” or “predicts failures.”
We detect silent model swaps and behaviour shifts against each agent’s own baseline, and show the numbers.	”Predicts model drift” or “prevents drift.”
Drift alerts are watch-only: a person decides.	”Automatically fixes drift.”
Sentinel runs on your own key, explains and recommends, and never enforces.	”Our AI agent stops your other AI agents” or “Sentinel catches attacks.”
Sentinel’s quality has not been measured yet.	Any accuracy, precision or “catches X percent” figure for Sentinel.
Bring your own monitor: your Claude, OpenAI, Gemini or script reviewer advises next to Sentinel; Sibyl’s policy decides.	”Your AI agent can now control your fleet” or “monitors stop bad actions.”
Monitors read metadata within the agent groups you choose.	”Your monitor sees everything” or “full visibility for your reviewer.”
Try the public demo at https://sibyl-demo.onrender.com, or a Codespaces link for your own copy.	”Free sandbox” or “trial of your own data” (the public demo is shared dummy data).
Native Anthropic and Gemini at the gateway; fallback stays within one protocol.	”Fails over from Claude to GPT automatically” for a native Anthropic

What we say	What we never say
	client.
One line or zero code for the listed clients and frameworks.	”Works with every framework” (Google ADK and AutoGen connectors are planned).
Learned models ship only if they beat the rules on your own labelled data.	”Learns from your data” or “gets smarter over time.”
We do not claim a data moat.	”Data moat” or “network intelligence.”
No LLM sits in the decision path.	”AI-powered governance.”
Sibyl owns routing, stopping, handoff and evidence; you own transaction integrity.	”Sibyl recovers your transactions” or “guarantees no duplicates.”
Supports human-oversight and record-keeping duties under named frameworks.	”EU AI Act compliant” or “certified.”
Measured gateway overhead is about 0.2 ms at p99 in our benchmark.	”Zero latency.”
Self-hosted; data stays in your cloud; no cross-tenant pooling.	”Bank-grade” or “military-grade security.”
Business-hours support, stated in writing.	”24/7 support.”
Enforcement covers traffic through the gateway and SDK.	”Covers every agent in your company.”

Claims to avoid

Until measured or built, never use: “predictive intelligence”, “learns from your data”, “data moat”, “prevents all unauthorised actions”, “predicts drift”, any quality figure for Sentinel or for an external monitor, “monitors enforce”, “public demo”, “AI that governs your AI”, “compliant” or “certified” for any framework, “guaranteed recovery”, “SOC 2”, “24/7”, “residency guaranteed” (residency is a declared field until region-pinned stacks exist), and any competitor feature claim marked **[unverified]** in chapter 5.

Glossary

Term	Meaning
Action Gate	Pre-execution allow, block or escalate on a proposed tool call, from signed policy
Action policy	Ordered rules on tool, arguments, cumulative caps and data class
Agent	An AI workflow Sibyl governs
Anchor	A signed checkpoint of the audit chain
Break-glass key	An emergency key that can only stop agents, usable when SSO is down
Continuity plan	Recovery targets and what to do on failure
Discovered agent	An agent seen in telemetry but not yet confirmed; never routed until an admin confirms it
Drift alert	`model_change` (a new served model or fingerprint) or `silent_drift` (a behaviour shift against the baseline); watch-only
External monitor	The customer's own agent or script registered to watch some or all agent groups and post advisory findings; never enforces
Monitor key	A key bound to one monitor that can only read metadata and post findings
Fallback chain	Ordered alternatives: model, provider, backup agent, human queue
Four-eyes	Two different people must approve
Governed autonomy	Share of production agent actions under an approved policy at an enforcement point
Handoff packet	Structured hand-over of unfinished work to a person
Held action	An escalated action waiting for a human decision; defaults to deny on timeout
Hold	A 30-minute pause on automation after an undo
Run tree	One multi-agent run: agents, steps, model and tool calls, handoffs

Term	Meaning
Sentinel	Optional advisory agent on the customer’s own model key; explains and recommends, never enforces
Tier	Business criticality, 1 (critical) to 3 (routine)
Verdict	An operator’s judgement on an alert: true positive, false positive or unsure

Addendum: plug-and-play monitoring, drift detection and Sibyl Sentinel (built)

Status for sales, 2026-10-12.

Plug-and-play monitoring

What to say. “Add one line, or run your app with `sibyl-run`, and Sibyl picks up the OpenAI, Anthropic and Google Gen AI clients and LangChain/LangGraph, the OpenAI Agents SDK and CrewAI. Multi-agent systems show up as a run tree. Agents we have never seen appear as discovered and are not routed until you confirm them.”

- Tools run by those frameworks go through the Action Gate before they execute.
- Prompts and outputs leave the process only if the capture mode allows, after the secrets scrub.
- **Limits to state.** Python only (TypeScript SDK planned). Google ADK and AutoGen connectors are planned.

Native Anthropic and Gemini at the gateway

What to say. “Your Anthropic or Gemini SDK changes only its base URL and key. Kill switch, limits, routing and the Action Gate work the same way.”

- **Limit to state.** Fallback stays inside one protocol. An agent on the native Anthropic protocol fails over to Anthropic-protocol targets only; Sibyl does not translate requests between providers. An agent that needs cross-provider failover uses the OpenAI-compatible protocol.

Model drift detection

What to say. “Sibyl tells you when a provider serves a different model than before, and when an agent’s behaviour shifts against its own baseline: output length, latency, refusals, truncation, tool use, schema failures, errors and tool sequences. Each alert

shows the numbers, including a noise floor so small samples do not raise alarms."

- It is plain statistics (PSI and Jensen–Shannon), no LLM.
- It is watch-only: it alerts a person, and it never fails over, holds an action or changes a policy on its own.
- **We detect; we do not predict.** Never say "predicts drift".
- **Limits to state.** Thresholds are tunable per tenant through the API (no console screen yet). A baseline needs at least 500 calls over 6 hours before it is compared. Fingerprint checks need a provider that sends one. No harness result yet.

Sibyl Sentinel

Prospects will ask for "an agent that watches our agents" built on Claude, GPT or Gemini. Our answer is **Sibyl Sentinel**: an optional, opt-in advisory agent. It runs on the customer's own Anthropic, OpenAI or Gemini key, and it explains incidents, reviews held actions for the approver, flags purpose drift and injection-like sequences, and writes a daily fleet summary.

How to position it.

- **Bring your own key.** The key is write-only, sealed in the tenant vault and never shown again. Calls go only to the provider's API host.
- **Advisory.** It cannot allow, block, approve, fail over or change policy. Allow, block, approve and failover stay deterministic, signed and auditable.
- **Bounded.** A daily budget the customer sets, reserved before each call so it is not overshot, and a sampling rate for drift reviews.
- **Respects capture mode.** Metadata-only agents send no prompts or outputs.
- **Accountable.** Every finding records provider, model, tokens, cost and a prompt hash. People rate findings as useful, not useful or wrong.
- Sell it as judgement and explanation on top of guarantees, not as the guarantee.
- If the prospect would rather use their own reviewer agent, that is now built too: see "Bring your own monitor agents" in chapter 4. Sentinel and external monitors can run side by side.

What we say. "Sentinel uses your preferred model to explain and recommend; Sibyl's enforcement never depends on a model's opinion."

What we never say. “Our AI agent stops your other AI agents.” The Action Gate stops actions; Sentinel advises. Never quote a quality figure: Sentinel has not been measured on the harness yet.

Action Gate status for sales

The policy evaluator, enforcement at the gateway (all three protocols) and in the SDK (including framework tools), and the approval workflow are built and tested: held actions, four-eyes approval, deny on timeout, cumulative caps, on-behalf-of with self-approval refused, and a decision history on every hold. The console screens run on the live API: the held-actions queue, the policy editor with simulate, and decision records in the audit log. In the merged harness runs, no unauthorised action executed in the Action Gate scenarios (S-29 at the SDK, S-30 at the gateway). The full gate suite is not yet green: every scenario still fails a known contract check (B-04), and the human-approval path is to be re-run after the B-09 fix, so we cite the scenario result, not a release gate. Known gaps to mention if asked: the held-actions list shows no arguments (the detail view shows them, scrubbed), and simulate does not read live counters.